Service Deployment and Timeouts

This guide documents default timeout and worker configuration for services deployed behind our API gateway.

The gateway enforces a hard 30 second upstream timeout. Services should configure their own internal timeouts well below that, typically 10-15 seconds, so a service can return its own error response instead of the gateway returning a generic 504.

Worker processes are sized based on expected concurrency: CPU-bound services run one worker per core, while I/O-bound services (most API services) run 2-4 workers per core since workers spend most of their time waiting on network calls.